

Search-Based Inference of Polynomial Metamorphic Relations

Jie Zhang, Junjie Chen, Dan Hao, Yingfei Xiong, Bing Xie*, Lu Zhang, Hong Mei
Key Laboratory of High Confidence Software Technologies (Peking University), MoE, China
Institute of Software, School of EECS, Peking University, China
{zhangjie12, chenjj14, haod, xiongyf04, xiebing, zhanglu, meih}@sei.pku.edu.cn

ABSTRACT

Metamorphic testing (MT) is an effective methodology for testing those so-called “non-testable” programs (e.g., scientific programs), where it is sometimes very difficult for testers to know whether the outputs are correct. In metamorphic testing, metamorphic relations (MRs) (which specify how particular changes to the input of the program under test would change the output) play an essential role. However, testers may typically have to obtain MRs manually.

In this paper, we propose a search-based approach to automatic inference of polynomial MRs for a program under test. In particular, we use a set of parameters to represent a particular class of MRs, which we refer to as polynomial MRs, and turn the problem of inferring MRs into a problem of searching for suitable values of the parameters. We then dynamically analyze multiple executions of the program, and use particle swarm optimization to solve the search problem. To improve the quality of inferred MRs, we further use MR filtering to remove some inferred MRs.

We also conducted three empirical studies to evaluate our approach using four scientific libraries (including 189 scientific functions). From our empirical results, our approach is able to infer many high-quality MRs in acceptable time (i.e., from 9.87 seconds to 1231.16 seconds), which are effective in detecting faults with no false detection.

Categories and Subject Descriptors

D.2.5 [Software Engineering]: Testing and debugging—Testing tools

General Terms

Algorithm, Experimentation, Verification

Keywords

Metamorphic testing; Invariant inference; Particle swarm optimization

*Corresponding author

1. INTRODUCTION

In software testing, a test oracle is needed to determine whether the program under test exhibits an acceptable behavior. Typically, it is very costly for testers to obtain a suitable test oracle for the program under test [28]. For some programs (e.g., scientific programs), which Weyuker [66] refers to as “non-testable” programs, obtaining test oracles may be extremely difficult or even impossible. In such a circumstance, testers may be unable to decide whether the program outputs are correct for most given inputs.

To facilitate testing the so-called “non-testable” programs, Chen et al. [11] proposed metamorphic testing (MT), which detects program faults in by looking for violations of metamorphic relations (MRs). Typically, an MR specifies how a particular change to the input would change the output. For example, to test the *sin* program, in metamorphic testing, the tester can check whether $\sin(x + \pi)$ is equal to $-\sin(x)$ without knowing the exact value of $\sin(x)$. That is, the MR “ $\sin(x + \pi) = -\sin(x)$ ” can be used to test the *sin* function.

Typically, testers need to manually acquire MRs. However, many MRs may be difficult to acquire due to the lack of thorough knowledge of the program under test. As scientific programs usually deal with complex computation, it is even more difficult to manually acquire MRs for scientific computing without thorough knowledge of such programs. For example, for the *sin* program, some MRs like “ $\sin(x + \pi) = -\sin(x)$ ” are easy to infer but some MRs like “ $\sin^2(\pi/2 - x) + \sin^2(x) = 1$ ” are not. As demonstrated by Chen et al. [14], more MRs may help achieve more adequate testing. Such manual acquisition of a number of MRs may become a bottleneck of metamorphic testing.

To facilitate this labor-intensive work, Kanewala and Bie-man [33] proposed a machine-learning based technique to predict whether a program contains some forms of MRs. For example, for the *sin* program, this technique predicts whether it contains a form of MR like “ $\sin(x) - c_1 \sin(x + c_2) = 0$ ” where c_1 and c_2 are constants, but does not generate the values of constants c_1 and c_2 . Different from their technique, in this paper, we aim to automatically infer specific MRs instead.

As it is difficult to infer arbitrary relations automatically, we need to consider a particular class of relations. After manually studying 70 MRs reported in the literature [38, 12, 47, 68, 83, 14, 70, 58, 59, 71, 44], we found that 43 (61.4%) MRs are polynomial, i.e., the relations between inputs and the relations between outputs are both polynomial equations. Therefore, in this paper, we focus on inferring only such polynomial MRs.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

ASE’14, September 15–19, 2014, Sweden

Copyright 2014 ACM 978-1-4503-3013-8/14/09 ...\$15.00.

http://dx.doi.org/10.1145/2642937.2642994

In particular, we propose a novel search-based approach to inferring polynomial MRs¹ by analyzing multiple executions of the program under test. Our approach aims to infer polynomial metamorphic relations in particular forms, which can be represented as a series of parameters. For those parameterized MRs, our approach uses particle swarm optimization [54] to search for suitable values of the parameters. That is, our approach tries to find a set of parameter values that characterize the analyzed executions in the form of an MR. Furthermore, to improve the quality of MRs, our approach applies statistics based filtering to the inferred MRs. Based on our study on the literature of MRs [38, 47, 68, 83, 12, 14, 70, 58, 59, 71, 44], 23 (53.5%) polynomial MRs are simple polynomial MRs whose relations between inputs are linear equations and whose relations between outputs are linear or quadratic equations. Therefore, in the current stage our approach focuses on inferring these particular polynomial MRs. However, our approach is general and can be used to infer other polynomial MRs, which is illustrated in Section 7.2 as an extension to our approach.

Based on our approach, we implemented a tool named MRI (i.e., Metamorphic Relation Inferer) and conducted three empirical studies on MRI to evaluate our approach using 189 scientific functions from Apache, JDK, GSL, and MATLAB. The first study investigates the feasibility of our approach. The second study investigates the quality of MRs inferred by our approach. The third study investigates whether MR filtering in our approach may improve the quality of inferred MRs. Our empirical results demonstrate that our approach is able to infer many high-quality MRs in acceptable time. Furthermore, MR filtering in our approach actually improves the quality of inferred MRs.

In summary, this paper makes the following contributions.

- A novel search-based² approach to inferring polynomial MRs via dynamically analyzing multiple executions using particle swarm optimization.
- An evaluation on four scientific libraries demonstrating the feasibility of our approach and the quality of inferred MRs.

2. EXAMPLES

To further motivate our research, we use some example metamorphic relations for three trigonometric functions to demonstrate that even for some widely known functions there may be some unfamiliar metamorphic relations.

To most people, the \sin function is one of the most popular trigonometric functions. Let us consider the metamorphic relation “ $\sin^2(\pi/2 - x) + \sin^2(x) = 1$ ”. At a glance, this metamorphic relation looks quite unfamiliar, but as we know that “ $\cos^2(x) + \sin^2(x) = 1$ ” and “ $\cos(x) = \sin(\pi/2 - x)$ ”, we may derive this metamorphic relation by ourselves. For the \cos function, “ $\cos(2x) = 2\cos^2(x) - 1$ ” is a metamorphic relation. Deriving this metamorphic relation would be much more difficult, because it involves complex mathematical deductions. Furthermore, it may be even more difficult to derive the metamorphic relation “ $\tan^2(x) - 2\tan(x)\tan(2x -$

$3\pi/2) - 1 = 0$ ” for the \tan function, as \tan is much less used than \sin and \cos .

As mathematicians have thoroughly investigated all the preceding trigonometric functions, testers may find the preceding metamorphic relations by searching in a textbook or the Internet. However, when the program under test does not exactly match some thoroughly investigated function, it may be painful for testers to derive its metamorphic relations.

Previously, Kanewala and Bieman [33] proposed a machine learning based approach to predicting likely metamorphic relations by some features, which are extracted from the control flow graph of a program. However, this approach can only predict whether the program has a particular metamorphic relation or not. Let us take the \sin function as an example. By analyzing the source code implementing the \sin function, this machine-learning based approach may predict that this function has an MR whose input change and output change are both additive³. That is, for the \sin function, this machine-learning based approach may infer metamorphic relations like $\sin(x) + c_1\sin(x + c_2) = 0$, where c_1 and c_2 are constants and x is the input parameter of the function \sin . However, this approach does not provide the values of constants c_1 and c_2 at all.

3. APPROACH

Before presenting our approach, we present some background information of particle swarm optimization (PSO) in Section 3.1. In our approach, we present definitions of MRs in their parameterized forms in Section 3.2, our PSO-based search algorithm for determining the parameters in Section 3.3, and our statistical based technique for filtering MRs in Section 3.4.

3.1 Background

Particle swarm optimization (PSO) [54], originally proposed by Kennedy and Eberhart [34] in 1995, is a swarm intelligence optimization algorithm simulating the birds foraging behavior. In PSO, each candidate solution is called a particle, and multiple particles coexist and optimize cooperatively to achieve the optimal solution. In particular, each particle has a velocity and a location, which keep changing during the search. There is a fitness function to evaluate how close the location of a particle is to an optimal location.

Consider searching in a D -dimensional space with N particles. We use $V_i^t = \langle v_{i1}^t, v_{i2}^t, \dots, v_{iD}^t \rangle$ and $L_i^t = \langle l_{i1}^t, l_{i2}^t, \dots, l_{iD}^t \rangle$ to denote the velocity and the location of the i -th ($1 \leq i \leq N$) particle at moment t ($t=1, 2, \dots$), where v_{id}^t and l_{id}^t denote the values of V_i^t and L_i^t in the d -th ($1 \leq d \leq D$) dimension, respectively. Then Formula 1 calculates the velocity of the i -th particle in the next moment (i.e., moment $t+1$).

$$v_{id}^{t+1} = \omega v_{id}^t + \xi_1 r_1 (p_{id}^t - l_{id}^t) + \xi_2 r_2 (p_{gd}^t - l_{id}^t) \quad (1)$$

In Formula 1, ω , ξ_1 , and ξ_2 are three weights in positive numbers (where ω is referred to as the inertia weight and the other two weights referred to as the acceleration factors); r_1 and r_2 are two random numbers between 0 and 1; p_{id}^t is the d -th dimension of the personal optimum location that the

¹Strictly speaking, the MRs our approach inferred are likely MRs.

²Search based software engineering [26], is an important branch of software engineering, which applies search techniques to solve various software-engineering problems [10, 8, 43]. To our knowledge, our approach is the first application of search based software engineering in MR inference.

³If the input to a program is modified by adding or subtracting a constant, its output will remain or increase. Kanewala and Bieman [33] defined such a change to an input/output as additive.

i -th particle has reached on and before moment t ; and p_{gd}^t is the d -th dimension of the global optimum location that all particles have reached on and before moment t . In other words, besides the two random numbers r_1 and r_2 , there are three factors that may impact the velocity of a particle in moment $t + 1$: the velocity in moment t , the personal optimal location that the particle has ever reached, and the global optimal location that any particle has ever reached. Based on its velocity in moment $t + 1$, the location of the i -th particle can be calculated with Formula 2.

$$l_{id}^{t+1} = l_{id}^t + v_{id}^{t+1} \quad (2)$$

In the beginning (i.e., moment 1), the N particles are assigned with locations and velocities randomly. The N particles keep updating their velocities and locations according to Formulae 1 and 2 until reaching moment T (which is a termination threshold). Then, the global optimum location that any particle has ever reached is returned as the final solution. Typically, each dimension is given a range, and if the location of any particle at any moment during the search is out of range, the location value in the corresponding dimension is set to the boundary values.

3.2 Parameterizing MRs

As a metamorphic relation (MR) usually specifies how a change to the input would result in a change to the output [33], we formalize an MR⁴ as Formula 3.

$$\mathbb{R}_i(I_1, I_2) \Rightarrow \mathbb{R}_o(O_1, O_2) \quad (3)$$

In Formula 3, I_1 and I_2 are the original input and the changed input, O_1 and O_2 are the outputs corresponding to I_1 and I_2 , \mathbb{R}_i is a relation between input I_1 and input I_2 , and \mathbb{R}_o is a relation between output O_1 and output O_2 . Without loss of generality, either I_1 or I_2 represents a vector of values, but either O_1 or O_2 represents only one value. The reason is that we can treat each output value individually when the program under test has more than one output value.

Although Formula 3 can exactly characterize an MR, it still cannot be directly used for MR inference, because both \mathbb{R}_i and \mathbb{R}_o can be in any form. As most MRs studied in the literature are polynomial, we further confine \mathbb{R}_i and \mathbb{R}_o in Formula 3 to polynomial equations. With this confinement, an MR can be characterized with the values of the parameters in the two polynomial equations.

Furthermore, as mentioned in the introduction, most polynomial metamorphic relations are simple polynomial metamorphic relations whose relations between inputs are linear equations and whose relations between outputs are linear equations or quadratic equations. That is, many MRs discussed in the literature are typically in much simpler forms than MRs involving two arbitrary polynomial equations. Therefore, in the current stage of our research, there are two cases under investigation: (1) \mathbb{R}_i is a linear equation, and \mathbb{R}_o is a linear equation, too. (2) \mathbb{R}_i is a linear equation, while \mathbb{R}_o is a quadratic equation. We will give a formula denoting MRs with each case.

Suppose that there are n input values for each input with $\langle x_1, x_2, \dots, x_n \rangle$ and $\langle y_1, y_2, \dots, y_n \rangle$ representing the two in-

puts (i.e., I_1 and I_2). As we mentioned, I_1 is the original input and I_2 is the input derived from I_1 . As \mathbb{R}_i is a linear equation that represents the relation between I_1 and I_2 , every input value of I_2 (i.e., y_i) can be denoted as a linear combination of all input values of I_1 , i.e., $y_i = \sum_{j=1}^n a_{ij}x_j + b_i$.

When \mathbb{R}_o is a linear equation, we can represent \mathbb{R}_o by $c_1O_1 + c_2O_2 + d = 0$, where O_1 and O_2 are the two outputs. Given a program under test (denoted as P), let us use $P(x_1, x_2, \dots, x_n)$ to denote output O_1 , and $P(\sum_{j=1}^n a_{1j}x_j + b_1, \dots, \sum_{j=1}^n a_{nj}x_j + b_n)$ to denote output O_2 . Then \mathbb{R}_o can be further represented as:

$$c_1P(x_1, x_2, \dots, x_n) + c_2P(\sum_{j=1}^n a_{1j}x_j + b_1, \dots, \sum_{j=1}^n a_{nj}x_j + b_n) + d = 0 \quad (4)$$

Similarly, we can have a formula denoting MRs with \mathbb{R}_i to be a linear equation and \mathbb{R}_o to be quadratic. As \mathbb{R}_o is quadratic, we can represent \mathbb{R}_o by $c_1O_1^2 + c_2O_1O_2 + c_3O_2^2 + d_1O_1 + d_2O_2 + e = 0$, where O_1 and O_2 are the two outputs. For simplicity, we use $P(I_1)$ and $P(\alpha I_1 + \beta)$ to denote the two outputs O_1 and O_2 , where α represents the matrix $a[i, j]$ and β represents the vector $\langle b_1, b_2, \dots, b_n \rangle$. Formula 5 denotes MRs with \mathbb{R}_o to be quadratic.

$$c_1P^2(I_1) + c_2P(I_1)P(\alpha I_1 + \beta) + c_3P^2(\alpha I_1 + \beta) + d_1P(I_1) + d_2P(\alpha I_1 + \beta) + e = 0 \quad (5)$$

It should be noted that Formula 4 is actually a degenerate form of Formula 5, where the three coefficients of the three terms with degree 2 are all zero. However, in our search algorithm discussed in Section 3.3, we need to force the algorithm to produce non-degenerate MRs. For example, if we allow the algorithm to produce degenerate MRs, it may produce all zero for all the parameters. Therefore, we use both formulae in our search algorithm without allowing any degenerate MRs.

In the following, we demonstrate that all the MRs discussed previously can be represented with our two formulae. For “ $\sin(x + \pi) = -\sin(x)$ ”, when $a_{11} = 1$, $b_1 = \pi$, $c_1 = 1$, $c_2 = 1$, and $d = 0$, Formula 4 actually represents this MR. Using Formula 5, when $a_{11} = -1$, $b_1 = \pi/2$, $c_1 = c_3 = 1$, $c_2 = 0$, $d_1 = d_2 = 0$, and $e = -1$, it represents “ $\sin^2(\pi/2 - x) + \sin^2(x) = 1$ ”; when $a_{11} = 2$, $b_1 = 0$, $c_1 = 2$, $c_2 = c_3 = 0$, $d_1 = 0$, $d_2 = -1$, and $e = -1$, it represents “ $\cos(2x) = 2\cos^2(x) - 1$ ”. Similarly, we can also use Formula 5 to represent “ $\tan^2(x) - 2\tan(x)\tan(2x - 3\pi/2) - 1 = 0$ ”.

In fact, given a program under test, the number of MRs that can be represented in Formula 4 or Formula 5 might still be infinite. Thus, it might still be infeasible to infer all the MRs satisfying the two formulae.

For ease of presentation, we refer to an MR inferred by using Formula 4 a type-one MR (denoted as 1-MR), whereas an MR inferred by using Formula 5 is called a type-two MR (denoted as 2-MR).

3.3 Searching for MR Parameters

Given a program under test (denoted as P) and M inputs (denoted as I_1, I_2, \dots, I_M) of P , the problem of inferring a polynomial MR for P can be turned into a search problem of finding some vector of parameter values in Formula 4 or Formula 5 such that for (almost) every input I_i ($1 \leq i \leq M$)

⁴Chen et al. [14] further generalized an MR as “an expected relation among the inputs and outputs of multiple executions”. We will investigate these generalized MRs in our future work.

the vector of parameter values and I_i satisfy Formula 4 or Formula 5.

To solve the preceding search problem, we adopt a PSO algorithm due to the following reasons. First, as PSO is very effective to search in continuous space, it may help find parameter values in real numbers. Second, as the location updating mechanism in PSO can keep particles from swaying among multiple optimal locations, it is suitable for the situation that there may be many MRs satisfying Formula 4 and/or Formula 5. However, when there are not many MRs, PSO is also effective due to the global optimal location and the velocity of the previous iteration can lead a particle to escape local optimal locations.

For simplicity of presentation, we focus on Formula 4 in the rest of this subsection. It is straightforward to extend to Formula 5. In our PSO, a candidate solution (i.e., a particle) is a set of parameter values (i.e., values of c_1 , c_2 , a_{ij} , b_i and d). For every set of parameter values, the fitness function counts the number of inputs that satisfy Formula 4.

Formally, we define the fitness function for Formula 4 as follows. Given a vector (denoted as L) of values for c_1 , c_2 , a_{ij} ($1 \leq i, j \leq n$), b_i ($1 \leq i \leq n$), and d , if L and input I_k ($1 \leq k \leq M$) satisfy Formula 4, we define $f(L, k) = 1$; otherwise we define $f(L, k) = 0$. Thus, the fitness of vector L can be defined as Formula 6, which actually counts the number of inputs that satisfy Formula 4 for L .

$$fitness(L) = \sum_{k=1}^M f(L, k) \quad (6)$$

As the location of every particle keeps changing, sometimes the values of c_1 and c_2 may be close to zero, which will make our MR pointless. To prevent our PSO algorithm from producing MRs in degenerate forms, when both values of c_1 and c_2 in a particle are close to zero, we reset them to a new value. In particular, we use a threshold (denoted as φ). When both c_1 and c_2 are between $-\varphi$ and φ , we set them to $-\varphi$ or φ depending on their being positive or negative.

As one execution of our PSO algorithm generates only one possible MR, we need to execute our PSO algorithm many times to obtain a number of MRs. Due to the random initialization of locations and velocities of the particles and the random factors (i.e., r_1 and r_2) in Formula 1, different executions may produce different MRs. It is possible that an execution of our PSO algorithm may not always produce a good enough solution (whose fitness is lower than a threshold denoted as F). In such cases, we drop all not good enough solutions. Specific parameter-setting (e.g., F) of our PSO algorithm is presented in Section 4.

3.4 MR Filtering

Our PSO algorithm infers MRs based on the multiple executions of the program, and thus the quality of inferred MRs may be dependent on the test inputs of a program. Intuitively, our PSO algorithm tends to produce high-quality MRs if many test inputs are used. However, it is costly to conduct our PSO algorithm with a large number of test inputs. Therefore, our proposed approach uses statistics based filtering to remove low-quality MRs.

For each MRs inferred by our PSO algorithm, the statistics based filtering algorithm applies a large number of randomly generated test inputs to program P , and records whether this MR is violated by each test input. If an MR is violated by an unignorable percentage (which is denoted

as S) of test inputs, we deem such an MR as a low-quality MR and remove it from the set of inferred MRs. By repeating the preceding process several times (which is denoted as Nof), we deliver a set of high-quality MRs by removing some low-quality MRs.

4. IMPLEMENTATION

Based on the approach described in Section 3, we implemented a tool named MRI (Metamorphic Relation Inferer).

For our PSO algorithm, we use the following settings as recommended in the literature of PSO [60, 36, 61]: we set the two acceleration factors (i.e., ξ_1 and ξ_2 in Formula 1) as 1.49445, the number of particles (i.e., N) as 20, and the termination threshold (i.e., the total number of moments T) as 350. The inertia weight (i.e., ω in Formula 1) is a changing value for different moments. We set the value of ω for moment t (denoted as ω^t) according to Formula 7 (also recommended in PSO literature), where ω_s (which is the value for moment 1) is 0.9 and ω_e (which is the value for moment T) is 0.4.

$$\omega^t = \omega_s - (\omega_s - \omega_e)(t/T)^2 \quad (7)$$

We set φ (which is the threshold to avoid degenerate MRs) as 0.5 and F (which is the threshold to select good enough solutions) as $95\% \cdot M$. Note that, in Formula 4, parameters c_1 , c_2 , and d can be inflated. If we multiply the three parameters with a common factor, we can have another MR that is semantically equivalent to the original MR. Thus, there can be a large number of MRs with the values of c_1 , c_2 , and d close to zero. Therefore, using a value significantly larger than zero for φ would also help save the efforts of inferring too many semantically equivalent MRs. Of course, 0.5 may not be the best value to achieve this goal. The setting of F as $95\% \cdot M$ is due to statistical considerations, as such a setting implies that it is highly probable that an input can satisfy the inferred MR. Similarly, we also set S as 5%.

As PSO algorithm requires to set the boundary values for the parameters c_1 , c_2 , c_3 , d_1 , d_2 , α , β , and e , we conducted a trial using the *sin* function of Apache 2.2 to decide the proper boundary values of these parameters, with which MRI may infer high-quality MRs. In particular, we varied the boundary values for these parameters and applied MRI with any specified values of these parameters to the *sin* function. After manually checking their inferred MRs for the *sin* function, we determined the boundary values for these parameters based on the effectiveness of their inferred MRs. As a result, in our implementation, the boundary values for c_1 , c_2 , c_3 , d_1 , d_2 , and α are from -2 to 2, the boundary value for β is from -10 to 10, and the boundary value for e is from 0 to 10.

Moreover, in MRI, as the test inputs used in our PSO algorithm and MR filtering are randomly generated, we set the range of test inputs to be from 0 to 20. Besides, the number of test inputs used in our PSO algorithm is set to 100, the number of test inputs used in MR filtering is set to 1000, and Nof is set to 10.

5. EMPIRICAL SETUP

To evaluate our approach, we conducted three empirical studies on 189 scientific functions from four scientific libraries to answer the following research questions

In the first study, we investigate the feasibility of our approach. In this study, we are interested in the following research question (RQ1): Is our approach able to infer MRs?

In the second study, we investigate the quality of inferred MRs. In this study, we are interested in the following research question (RQ2): How is the quality of MRs inferred by our approach?

In the third study, we investigate the impact of MR filtering in our approach. In this study, we are interested in the following research question (RQ3): Does MR filtering improve the quality of inferred MRs?

5.1 Subjects

Our approach infers MRs by analyzing multiple executions of a program rather than its source code. That is, our approach is a black-box technique, which requires no source code under test and may be applied to scientific functions in various languages. Therefore, in our empirical studies we used four scientific libraries which are written in Java and C/C++.

The four scientific libraries are Apache Commons Mathematics Library, the math class of Java Development Kit, the GNU Scientific Library, and the scientific library of MATLAB. Apache Commons Mathematics Library⁵ (abbreviated as Apache in this paper) is a library of self-contained mathematical and statistical components in Java. The math class of Java Development Kit⁶ (abbreviated as JDK in this paper) provides methods for numeric operations in the Java platform. The GNU Scientific Library⁷ (abbreviated as GSL in this paper) is a numeric library of mathematical routines in C/C++ language. MATLAB⁸ is a powerful computing environment developed by MathWorks, which also contains a very large number of scientific functions. Among the four libraries, Apache and GSL are open-source scientific libraries, whereas JDK and MATLAB are commercial scientific libraries. That is, the source code of the latter two libraries is not available.

As each scientific library contains a very large number of scientific functions, it is impossible for us to run our approach on all the scientific functions due to time limit on experimentation, so we further selected the packages that perform plain mathematical computations (e.g., not matrix operations or statistical functions). In particular, we used the scientific programs in *FastMath.java* of Apache. For JDK and MATLAB, we used their same scientific programs as Apache. For GSL, we used the scientific programs in the “specfunc” directory. As a result, we got a dataset consisting of 56 scientific functions of Apache 2.2, 53 scientific functions of JDK 1.6, 55 scientific functions of GSL 1.8, and 25 scientific functions of MATLAB R2012b. Table 1 presents the basic information of these scientific functions, where the last three columns present the version information, the total number of scientific functions, and the total number of lines of code used in the studies. More information on these functions can be found in the project webpage: <http://infermrs.sourceforge.net/>. As the source code of JDK 1.6 and that of MATLAB R2012b is not available, we do not list its number of lines of code in this table.

⁵<http://commons.apache.org/proper/commons-math>

⁶<http://www.oracle.com/technetwork/java/javase/downloads/index.html>

⁷<http://www.gnu.org/software/gsl/>

⁸<http://www.mathworks.cn>

Table 1: Subjects

Library	Version	#Function	LOC
Apache	2.2	56	1,626
JDK	1.6	53	-
GSL	1.8	55	7,309
MATLAB	R2012b	25	-

5.2 Process

In this subsection, we present the details of how we conducted the three empirical studies.

5.2.1 Study I

In the first study, for each function, we repeated our PSO algorithm 500 times. After MR filtering MRI inferred a set of 1-MRs and 2-MRs for each subject. We also recorded the total time spent on MR inference for each subject.

5.2.2 Study II

First, we investigated the correctness of MRs inferred in the first study. As *sin*, *cos*, and *tan* are typical scientific functions, which are available through many mathematics books, we chose the three functions implemented in different libraries as representative subjects and manually checked the correctness of their inferred MRs as follows. For each of these inferred MRs, we looked through Wikipedia⁹ and a mathematics book [22] to check whether it is correct.

Some inferred MRs can be deduced by others. For example, $\sin(-x) - \sin(x - \pi) = 0$ can be deduced by $\sin(x) + \sin(-x) = 0$ and $\sin(x) + \sin(x - \pi) = 0$. In software testing, if developers have already used the latter two MRs, they may not run again the *sin* function with the former MR (i.e., $\sin(-x) - \sin(x - \pi) = 0$) because faults detected by the former MR may also be detected by the latter MRs. To acquire the set of MRs that cannot be deduced by each other, we define a set of representative MRs for any set of inferred MRs. In particular, for any given set of inferred MRs (i.e., 1-MRs or 2-MRs) of a subject, its set of representative MRs is the minimized subset of the given set, and each MR of the given set can be deduced by using one or more MRs in its representative set. For example, for a set of inferred MRs $\{\sin(-x) - \sin(x - \pi) = 0, \sin(x) + \sin(-x) = 0, \sin(x) + \sin(x - \pi) = 0\}$, its set of representative MRs is $\{\sin(x) + \sin(-x) = 0, \sin(x) + \sin(x - \pi) = 0\}$. In this study, for each trigonometric function, we manually summarized its set of representative 1-MRs and 2-MRs respectively. If multiple minimal subsets exist, we chose one that has the smallest size.

Besides, Apache 2.2, JDK 1.6, and MATLAB R2012b implement some common scientific functions. For each common scientific function (e.g., *exp*), we compared the number of MRs inferred from different libraries.

Second, we investigated the fault-detection capability of MRs inferred in the first study through regression testing [69, 78]. In particular, to simulate the regression testing in software evolution, we applied the MRs inferred from the correct version of a project on detecting faults in its subsequent version. Among the functions used in the empirical studies, 11 functions (i.e., *sin*, *cos*, *tan*, *log10*, *loglp*, *asinh*, *atan*, and four *abs* functions with various inputs) have changed from Apache 2.2 to Apache 3.2. The other functions used in the

⁹[http://en.wikipedia.org/wiki/Trigonometric/_functions/#Sine.2C/_cosine/_and/_tangent](http://en.wikipedia.org/wiki/Trigonometric_functions/#Sine.2C/_cosine/_and/_tangent)

empirical studies have no change during software evolution. Therefore, in the second study we used these 11 functions of Apache 3.2 to investigate the fault-detection capability of MRs inferred from the first study. As the functions of Apache have been widely used in practical software development, they hardly contain any faults. Therefore, for these functions we constructed faults using program mutation following procedure similar to prior work [46, 77]. The difference between versions shows the developers’ modification on the previous version, and thus we generated faults only in such difference so as to simulate most developers’ faults in software evolution. In particular, for each of the 11 functions of Apache 3.2, we applied MuClipse [62] to generate a number of mutants whose mutation operators¹⁰ occur only on the different statements between Apache 2.2 and Apache 3.2. Each mutant, which is the result of applying a mutation operator on the source code, is viewed as a faulty program in our second study.

If an MR is violated by a faulty program, we deem the MR detects the fault. However, if the MR is also violated by the original, unseeded program of Apache 3.2, we deem this MR is of low quality and the detection is a false detection. For each of the 11 functions of Apache 3.2, we randomly generated 1000 test inputs. Then, we ran these test inputs on both the generated faulty versions of Apache 3.2 and the original version of Apache 3.2 for each MR inferred from Apache 2.2. As our approach infers MRs that are supported by at least 95% inputs, the inferred MRs should be used in a statistical way of metamorphic testing, which is referred to as statistical metamorphic testing in this paper. In statistical metamorphic testing, only when the violation of an inferred MR become statistically non-trivial, we deem the program to be likely to contain faults. In particular, we consider an MR was violated when at least 5% of the test inputs were violated considering anomaly detection.

5.2.3 Study III

To learn whether MR filtering improves the quality of inferred MRs, we compared the quality of MRs inferred with MR filtering and the quality of MRs inferred without MR filtering in regression testing following the same procedure of the second study. In this study, for each faulty program of Apache 3.2, we used statistical metamorphic testing to evaluate the quality of these MRs, recording the number of true detections and the number of false detections. Finally, we compared the numbers of true detections and false detections between MRs inferred with MR filtering and MRs inferred without MR filtering.

5.3 Threats to Validity

The threat to internal validity lies in the implementation of our approach. To reduce the threat from implementing errors, the authors of this paper reviewed the source code after implementing the proposed approach.

The main threats to external validity lie in the subjects and faults. First, similar to prior work [14] in metamorphic testing, we used four scientific libraries consisting of small¹¹

¹⁰Mutation operators define some operations like statement deletion, statement replacement, and so on.

¹¹After manually studying the scientific functions used in the literature of software engineering [5, 63, 7, 23], we found that the scientific functions used in their evaluation are usually very small, which are usually smaller than 100 lines of code.

Table 2: Basic statistics on MR inference

Number of MRs for each scientific program						
Library	1-MRs			2-MRs		
	Avg.	Max.	Min.	Avg.	Max.	Min.
Apache	87.96	353	0	46.63	401	0
JDK	85.04	348	0	47.72	395	0
GSL	80.85	331	0	52.78	239	0
MATLAB	47.24	168	0	13.52	78	0
Execution time of MR inference for each scientific program						
Library	1-MRs (seconds)			2-MRs (seconds)		
	Avg.	Max.	Min.	Avg.	Max.	Min.
Apache	49.87	261.03	15.19	33.33	96.20	16.25
JDK	42.37	404.08	14.92	45.35	421.88	17.02
GSL	97.70	315.05	9.87	352.88	1231.16	18.15
MATLAB	123.99	202.41	67.53	247.10	474.29	88.01

scientific functions whose MRs may be manually checked. As our approach is a black-box technique that does not analyze the source code, whether the scientific functions are small or large does not affect the feasibility of our approach. However, the functionality of scientific functions has much impact on the feasibility of our approach because our approach infers MRs based on their executions. To reduce this threat, we used a large number of scientific programs from different libraries. Moreover, although our approach is evaluated based on scientific programs, the approach has no such restrictions and can be applied to any programs. To reduce this threat, we will evaluate our approach by other programs in the future. Second, the faults were generated by using a mutation tool because prior work [2] shows that such faults can be used in the empirical studies of software testing. As these faults may be not representative of real faults, we will conduct more empirical studies on more programs with real faults in the future. Furthermore, in the evaluation we used the implementation of our approach introduced in Section 4, whose value for each parameter is set based on the literature of PSO and our trial on the *sin* function. However, as our approach does not have any restrictions on the values of the parameters, in future work we will conduct empirical studies to evaluate the effectiveness of our approach with other values of these parameters.

6. RESULTS

6.1 RQ1: MR Inference

Table 2 presents the basic statistics on MR inference, including the number of inferred MRs for each scientific function and the execution time of our approach on inferring MRs for each scientific function. The complete results can be found in <http://infermrs.sourceforge.net/>. From this table, for each scientific function, the number of 1-MRs is from 0 to 353, whereas the number of 2-MRs is from 0 to 401. That is, our approach infers an unignorable number of 1-MRs and 2-MRs for scientific functions. The execution time of our approach on inferring MRs for each scientific function is from 9.87 seconds to 1231.16 seconds, which is acceptable. Therefore, our approach is able to infer many MRs quickly.

6.2 RQ2: Quality of Inferred MRs

6.2.1 Correctness

Table 3 presents typical MRs of the three trigonometric functions *sin*, *cos*, and *tan*. Our approach generates MRs

Table 3: MRs inferred from three trigonometric functions

Function	Library	1-MRs	2-MRs
<i>sin</i>	Apache	$\sin(x) - \sin(x - 2\pi) = 0$ $\sin(x) + \sin(x - \pi) = 0$ $\sin(x) + \sin(-x) = 0$ $\sin(x) - \sin(-x + \pi) = 0...$	$\sin^2(x) + \sin^2(-x + 0.5\pi) - 1 = 0$ $\sin^2(-0.5x - 0.75\pi) + 0.5\sin(x) - 0.5 = 0$ $\sin^2(x) + \sin^2(-x) + 2\sin(x)\sin(-x) = 0$ $\sin^2(x) + \sin^2(x - \pi) + 2\sin(x)\sin(x - \pi) = 0...$
	JDK	$\sin(x) - \sin(x - 2\pi) = 0$ $\sin(x) + \sin(x - \pi) = 0$ $\sin(x) + \sin(-x) = 0$ $\sin(x) + \sin(-x + 4\pi) = 0...$	$\sin^2(x) + \sin^2(x - 2.5\pi) - 1 = 0$ $\sin^2(x) + 0.5\sin(2x - 1.5\pi) - 0.5 = 0$ $\sin^2(x) + \sin^2(-x) + 2\sin(x)\sin(-x) = 0$ $\sin^2(x) + \sin^2(x - \pi) + 2\sin(x)\sin(x - \pi) = 0...$
	GSL	$\sin(x) - \sin(-x - \pi) = 0$ $\sin(x) + \sin(x + \pi) = 0$ $\sin(x) + \sin(-x) = 0$ $\sin(x) - \sin(-x - 3\pi) = 0...$	$\sin^2(x) - \sin^2(-x - 2\pi) - \sin(x) - \sin(-x - 2\pi) = 0$ $\sin^2(x) + 0.5\sin^2(-x) + 1.5\sin(x)\sin(-x) + 0.5\sin(x) + 0.5\sin(-x) = 0$ $\sin^2(x) - \sin(x)\sin(-x + \pi) + \sin(x) - \sin(-x + \pi) = 0$ $\sin^2(x) + \sin^2(-x) - 2\sin(x)\sin(-x) = 0...$
	MATLAB	$\sin(x) + \sin(-x) = 0$ $\sin(x) + \sin(x - \pi) = 0$ $\sin(x) + \sin(-x + 2\pi) = 0$ $\sin(x) - \sin(x - 2\pi) = 0...$	$\sin^2(x) + \sin^2(-x + 0.5\pi) - 1 = 0$ $\sin^2(x) + \sin^2(-x - \pi) - 2\sin(x)\sin(-x - \pi) = 0$ $\sin^2(x) + \sin(x)\sin(x - \pi) + 2\sin(x) + 2\sin(x - \pi) = 0$ $\sin^2(x) + 0.5\sin(2x + 0.5\pi) - 0.5 = 0...$
<i>cos</i>	Apache	$\cos(x) + \cos(-x - \pi) = 0$ $\cos(x) - \cos(x - 2\pi) = 0$ $\cos(x) + \cos(x - \pi) = 0$ $\cos(x) - \cos(-x) = 0...$	$\cos^2(x) + \cos^2(-x - 0.5\pi) - 1 = 0$ $\cos^2(-0.5x - 1.5\pi) + 0.5\cos(x) - 0.5 = 0$ $\cos^2(x) - \cos^2(x + \pi) = 0$ $\cos^2(x) + \cos^2(-x) - 2\cos(x)\cos(-x) = 0...$
	JDK	$\cos(x) + \cos(-x - \pi) = 0$ $\cos(x) - \cos(x - 2\pi) = 0$ $\cos(x) + \cos(x - \pi) = 0$ $\cos(x) - \cos(-x) = 0...$	$\cos^2(x) + \cos^2(-x + 0.5\pi) - 1 = 0$ $\cos^2(0.5x - 2\pi) - 0.5\cos(x) - 0.5 = 0$ $\cos^2(x) + \cos(x)\cos(-x + \pi) - \cos(x) - \cos(-x + \pi) = 0$ $\cos^2(x) - \cos(x)\cos(-x) = 0...$
	GSL	$\cos(x) - \cos(x + 2\pi) = 0$ $\cos(x) + \cos(x - \pi) = 0$ $\cos(x) - \cos(-x) = 0$ $\cos(x) - \cos(-x - 3\pi) = 0...$	$\cos^2(x) + 0.5\cos^2(x + \pi) + 1.5\cos(x)\cos(x + \pi) - 1.5\cos(x) - 1.5\cos(x + \pi) = 0$ $\cos^2(x) - \cos(x)\cos(-x + 2\pi) - \cos(x) + \cos(-x + 2\pi) = 0$ $\cos^2(0.5x) - 0.5\cos(x) - 0.5 = 0$ $\cos^2(x) + \cos^2(x - 2\pi) - 2\cos(x)\cos(x - 2\pi) = 0...$
	MATLAB	$\cos(x) - \cos(-x + 2\pi) = 0$ $\cos(x) - \cos(-x) = 0$ $\cos(x) + \cos(x - \pi) = 0$ $\cos(x) + \cos(x + \pi) = 0...$	$\cos^2(-0.5x - 1.5\pi) + 0.5\cos(x) - 0.5 = 0$ $\cos^2(x) + 3\cos^2(x - \pi) + 4\cos(x)\cos(x - \pi) = 0$ $\cos^2(x) + \cos^2(x + 0.5\pi) - 1 = 0$ $\cos^2(x) - \cos^2(-x - \pi) + \cos(x) + \cos(-x - \pi) = 0...$
<i>tan</i>	Apache	$\tan(x) + \tan(-x + \pi) = 0$ $\tan(x) - \tan(x - 2\pi) = 0$ $\tan(x) - \tan(x - \pi) = 0$ $\tan(x) - \tan(x + 2\pi) = 0...$	$\tan^2(0.5x + 1.25\pi) - 2\tan(x)\tan(0.5x + 1.25\pi) - 1 = 0$ $\tan^2(x) + \tan^2(-x) + 2\tan(x)\tan(-x) = 0$ $\tan^2(x) - \tan^2(x + \pi) - \tan(x) + \tan(x + \pi) = 0$ $\tan^2(x) + \tan^2(x + 3\pi) - 2\tan(x)\tan(x + 3\pi) = 0...$
	JDK	$\tan(x) + \tan(-x + \pi) = 0$ $\tan(x) + \tan(-x + 2\pi) = 0$ $\tan(x) - \tan(x - \pi) = 0$ $\tan(x) - \tan(x + \pi) = 0...$	$\tan^2(x) - 2\tan(x)\tan(2x - 1.5\pi) - 1 = 0$ $\tan^2(x) - \tan^2(-x) = 0$ $\tan^2(x) - \tan^2(x + \pi) + \tan(x) - \tan(x + \pi) = 0$ $\tan^2(x) + \tan^2(-x - 2\pi) + 2\tan(x)\tan(-x - 2\pi) = 0...$
	MATLAB	$\tan(x) + \tan(-x) = 0$ $\tan(x) - \tan(x - 2\pi) = 0$ $\tan(x) + \tan(-x + \pi) = 0$ $\tan(x) + \tan(-x - 2\pi) = 0...$	$\tan^2(x) + \tan^2(-x + \pi) + 2\tan(x)\tan(-x + \pi) = 0$ $\tan^2(x) + 0.5\tan^2(-x) + 1.5\tan(x)\tan(-x) - 0.5\tan(x) - 0.5\tan(-x) = 0$ $\tan^2(x) + \tan^2(-x + 2\pi) + 2\tan(x)\tan(-x + 2\pi) = 0$ $\tan^2(0.5x - 1.75\pi) - 2\tan(x)\tan(0.5x - 1.75\pi) - 1 = 0...$

Table 4: Comparison of some inferred MRs in the first study

1-MRs	<i>abs_d</i>	<i>abs_f</i>	<i>abs_i</i>	<i>abs_l</i>	<i>acos</i>	<i>acosh</i>	<i>asin</i>	<i>asinh</i>	<i>atan</i>	<i>round_d</i>	<i>cbrt</i>	<i>ceil</i>	<i>cos</i>
Apache	141	160	64	73	0	19	0	21	62	118	23	123	207
JDK	132	145	96	72	0	-	0	-	60	120	30	97	213
MATLAB	148	-	-	-	19	8	44	10	50	78	-	12	129
1-MRs	<i>toDegrees</i>	<i>signum_f</i>	<i>expm1</i>	<i>floor</i>	<i>gE_d</i>	<i>gE_f</i>	<i>log</i>	<i>log10</i>	<i>log1p</i>	<i>nextUp_d</i>	<i>nextUp_f</i>	<i>rint</i>	<i>atanh</i>
Apache	0	302	73	115	106	106	40	133	50	331	350	126	0
JDK	5	306	70	120	115	117	60	135	54	348	334	110	-
MATLAB	-	-	-	84	-	-	28	95	44	-	-	-	52
1-MRs	<i>round_f</i>	<i>signum_d</i>	<i>exp</i>	<i>sin</i>	<i>sinh</i>	<i>sqrt</i>	<i>tan</i>	<i>tanh</i>	<i>cosh</i>	<i>toRadians</i>	<i>ulp_d</i>	<i>ulp_f</i>	
Apache	122	344	64	219	0	9	20	139	0	353	220	236	
JDK	122	317	68	227	0	13	13	134	0	336	217	229	
MATLAB	-	-	41	168	0	5	8	125	0	-	-	-	
2-MRs	<i>abs_d</i>	<i>abs_f</i>	<i>abs_i</i>	<i>abs_l</i>	<i>acos</i>	<i>acosh</i>	<i>asin</i>	<i>asinh</i>	<i>atan</i>	<i>round_d</i>	<i>cbrt</i>	<i>ceil</i>	<i>cos</i>
Apache	5	2	44	39	0	3	0	2	14	46	1	28	108
JDK	1	3	30	42	0	-	0	-	7	47	4	31	101
MATLAB	6	-	-	-	3	1	4	1	6	21	-	20	62
2-MRs	<i>toDegrees</i>	<i>signum_f</i>	<i>expm1</i>	<i>floor</i>	<i>gE_d</i>	<i>gE_f</i>	<i>log</i>	<i>log10</i>	<i>log1p</i>	<i>nextUp_d</i>	<i>nextUp_f</i>	<i>rint</i>	<i>atanh</i>
Apache	0	401	26	58	42	56	4	18	6	5	5	48	0
JDK	0	395	30	47	51	57	2	15	1	8	3	50	-
MATLAB	-	-	-	24	-	-	2	10	4	-	-	-	6
2-MRs	<i>round_f</i>	<i>signum_d</i>	<i>exp</i>	<i>sin</i>	<i>sinh</i>	<i>sqrt</i>	<i>tan</i>	<i>tanh</i>	<i>cosh</i>	<i>toRadians</i>	<i>ulp_d</i>	<i>ulp_f</i>	
Apache	38	391	31	131	0	0	8	80	0	303	214	246	
JDK	47	394	26	109	0	2	8	73	0	299	218	225	
MATLAB	-	-	9	78	0	7	6	58	0	-	-	-	

by assigning values to the parameters in the formulae with some precision. To ease understanding, we present the inferred MRs in this table by using the estimation of these values. For example, in the *sin* function, we use π to denote 3.14¹². GSL does not implement the *tan* function and thus we do not list its inferred MRs in this table. For each subject, we use the bold font to depict its complete sets of representative MRs. The left columns give 1-MRs inferred by our approach using Formula 4 whereas the right columns give 2-MRs inferred by our approach using Formula 5.

From this table, these typical MRs include most important MRs of the three trigonometric functions. For example, our approach infers series of 1-MRs and 2-MRs for the *cos* function of the four scientific libraries. These inferred 1-MRs, represented by $\cos(x) - \cos(-x) = 0$ and $\cos(x) + \cos(x - \pi) = 0$, show that the *cos* function is a symmetric and periodical function. The inferred 2-MRs, represented by $\cos^2(x) + \cos^2(-x - 0.5\pi) - 1 = 0$, $\cos^2(-0.5x - 1.5\pi) + 0.5\cos(x) - 0.5 = 0$, $\cos^2(x) - \cos^2(x + \pi) = 0$, and $\cos^2(x) + \cos^2(-x) - 2\cos(x)\cos(-x) = 0$, show the relation between $\cos(x)$, $\cos(2x)$ and $\cos(x - 0.5\pi)$ besides the symmetric and periodical characteristics of the *cos* function. Furthermore, our approach infers the complex MRs $\sin^2(\pi/2 - x) + \sin^2(x) - 1 = 0$ for the *sin* function and $\tan^2(x) - 2\tan(2x - 3\pi/2)\tan(x) - 1 = 0$ for the *tan* function from the libraries.

Table 4 presents the total number of MRs inferred from the common scientific functions. For each function, the numbers of inferred MRs from different libraries are close. This observation is as expected because these functions have the same functionality, suggesting the correctness of the inferred MRs. For the same scientific function, the average execution time of our approach for different libraries is close. That is, although different libraries may implement a scientific function in different ways¹³ (i.e., resulting in different programs), their kernel source code may not differ much in efficiency and thus the execution time of our approach for the same scientific functions of different libraries is close.

From Table 4, the number of inferred MRs is larger than that of the representative MRs. For example, our approach generated 219 1-MRs for the *sin* function of Apache, but only 2 of them are representative. The other 217 MRs can be deduced by these 2 representative MRs. As Chen et al. [11] demonstrate that more MRs may help to achieve more adequate testing, the MRs that can be deduced by some representative MRs may not be redundant. Moreover, more MRs may reduce the cost in software testing. For example, in order to reveal the faults that can be detected only by $\sin(x) - \sin(-x + 3\pi) = 0$, it may be more costly to check the two representative MRs (i.e., $\sin(x) + \sin(-x) = 0$ and $\sin(x) + \sin(x - \pi) = 0$) rather than one MR. To check the former MR, developers may run the *sin* function twice, whereas checking the latter two MRs, developers may run the *sin* function four times.

6.2.2 Fault-Detection Capability

Table 5 gives the results of the second study, where the second column gives the total number of mutation faults

Table 5: Fault-detection capability of MRs

	Seeded Faults	#By 1-MRs			#By 2-MRs		
		Total	FD	TD	Total	FD	TD
<i>sin</i>	17	9	0	9	9	0	9
<i>cos</i>	19	8	0	8	8	0	8
<i>tan</i>	18	8	0	8	8	0	8
<i>log10</i>	58	7	0	7	4	0	4
<i>log1p</i>	115	25	0	25	24	0	24
<i>asinh</i>	297	1	0	1	0	0	0
<i>atan</i>	94	15	0	15	31	0	31
<i>abs_d</i>	7	5	0	5	5	0	5
<i>abs_f</i>	7	5	0	5	5	0	5
<i>abs_i</i>	15	15	0	15	15	0	15
<i>abs_l</i>	15	15	0	15	15	0	15

in each scientific function of Apache 3.2, “FD” presents the number of false detections by the corresponding MRs, “TD” presents the number of true detections by the corresponding MRs, and “Total” is the sum of its previous two columns.

From the fifth and eighth columns, the numbers of true detections for 1-MRs and 2-MRs are usually larger than 0. That is, the inferred MRs are able to detect faults. From the fourth and seventh columns, the numbers of false detections for 1-MRs and 1-MRs are 0. That is, the inferred MRs always make correct detection. On the other hand, for most scientific functions (including *sin*, *cos*, *tan*, *abs_d*, *abs_f*, *abs_i*, and *abs_l*), the inferred MRs detect about half of the faults. The only exception is *asinh*, in which only one fault is detected out of 297 faults. By further investigating the injected faults, we found that this is probably because the rest 296 faults were not triggered by the test inputs: the mutated statements of all the 296 faults will be executed only when a strict condition (i.e., variable *a* is smaller than 0.167) is satisfied. Overall, our inferred MRs were effective in detecting faults and produced no false detection.

Comparing the results of the two types of MRs (i.e., 1-MRs and 2-MRs), the number of faults detected by the former is close to that of the latter. After reviewing these inferred MRs and their detected faults, we found the reason to be that our approach always infers some important MRs that can detect a large number of faults no matter which formula (i.e., Formula 4 and Formula 5) it used.

6.3 RQ3: Necessity of MR Filtering

Table 6 presents the fault-detection capability of MRs inferred by our approach without MR filtering. For each scientific function, the number of false detections of MRs inferred without MR filtering is usually larger than 0. For example, the 1-MRs for the *tan* function have 10 false detections. As the functionality of these subjects does not change from Apache 2.2 to Apache 3.2, these inferred MRs should not be violated. That is, the false detections in this table result from low-quality MRs, which are inferred by our approach without MR filtering. As the number of false detections detected by the MRs inferred with MR filtering is 0 (shown by Table 5), MR filtering actually improves the quality of MRs by removing low-quality MRs.

On the other hand, the filtering did reduce the number of true detections, but the reduced number was small. Overall, the reduction on fault detection occurred only on 3 out of 11 functions, and in total only 16.9% true detections were filtered out. Considering the large number of false detections removed, we believe the filtering procedure is effective and necessary.

¹²The complete list of inferred MRs for these functions can be found in the project webpage.

¹³As the source code of JDK and MATLAB is not available, we cannot check the difference between the source code of the three libraries.

Table 6: Fault-detection capability of MRs without MR filtering

	Seeded Faults	#By 1-MRs			#By 2-MRs		
		Total	FD	TD	Total	FD	TD
<i>sin</i>	17	9	0	9	9	0	9
<i>cos</i>	19	9	0	9	19	11	8
<i>tan</i>	18	18	10	8	18	10	8
<i>log10</i>	58	58	51	7	58	54	4
<i>log1p</i>	115	29	0	29	113	85	28
<i>asinh</i>	297	297	292	5	4	0	4
<i>atan</i>	94	94	65	29	94	63	31
<i>abs_d</i>	7	7	2	5	5	0	5
<i>abs_f</i>	7	7	2	5	7	2	5
<i>abs_i</i>	15	15	0	15	15	0	15
<i>abs_l</i>	15	15	0	15	15	0	15




Figure 1: A 3-dimension solution space

7. DISCUSSION

7.1 Limitations

First, our approach requires the input of the program under test to be of numerical values. For example, our approach in the current stage may not be suitable for programs whose inputs are pointers or arrays. Therefore, we may improve the existing PSO algorithms by transforming these inputs into numerical values.

Second, our approach requires that the program under test should always produce the same output for the same input. That is to say, the behavior of the program under test should not depend on some external state. Thus, our approach may not be suitable for testing a program involving a database or a method in an object-oriented program relying on the state of the object.

7.2 Extensions

According to Chen et al. [14], an MR are supposed to hold among multiple executions. Based on this, we can change our Formula 3 to a more general form depicted in Formula 8.

$$\mathbb{R}_i(I_1, I_2, \dots, I_m) \Rightarrow \mathbb{R}_o(O_1, O_2, \dots, O_m) \quad (8)$$

Similar to the treatment presented in Section 3, we can also parameterize this generalized definition of MRs by confining \mathbb{R}_i and \mathbb{R}_o to be polynomial equations. In fact, we can further relax the requirement of polynomial equations to polynomial inequalities.

Therefore, the whole solution space of the extension can be depicted in Figure 1, where P_1 and P_2 represent the two cases our approach has solved. In this solution space, there are three major directions for further extending our approach.

First, the \mathbb{R}_i relation on inputs can be extended to a polynomial with a higher degree (i.e., greater than 1). For example, using a polynomial with a degree of 2 for \mathbb{R}_i may help infer " $\log_{10}(x^2) = 2\log_{10}(x)$ " for the \log_{10} function.

Second, the \mathbb{R}_o relation on outputs can also be extended to a polynomial with a higher degree (i.e., greater than 2). For example, using a polynomial with a degree of 3 for \mathbb{R}_o may help infer " $\sin(3x) = 3\sin(x) - 4\sin^3(x)$ " for \sin .

Third, the number of involved inputs can be extended to more than 2. For example, involving 3 inputs may help infer " $\sin(2x) = 2\sin(x)\sin(\pi/2 - x)$ " for the \sin function.

As our approach in this paper cannot deal with these extensions, we will improve our PSO algorithm for further investigating these situations.

8. RELATED WORK

8.1 Metamorphic Testing

Chen et al. originally proposed the methodology of metamorphic testing [11] and formally established the methodology [15]. Besides application of metamorphic testing on various areas [12, 48], some researchers focus on the selection of good MRs, which is related to our work. Chen et al. [13] demonstrated that it is better to select MRs that make the multiple executions of the program as different as possible. Mayer and Guderlei [45] identified some MRs and classified them using mutation analysis. They also evaluated MRs according to their potential usefulness. Recently, Liu et al. [38] proposed to systematically construct MRs based on some already identified MRs. Differently, our approach aims to automatically infer MRs without any existing MRs.

Our work is mostly related to the technique proposed by Kanewala and Bieman [33], which automatically predicts the existence of some forms of MRs for a program using machine learning. Different from our approach, their technique does not produce specific MRs, but tells whether a program may have a particular form of MRs or not. Their technique and our approach may be viewed as complement to each other. In particular, we may use their technique to predict the existence of a form of MR and use our technique to produce the specific MR by giving the values of parameters. Furthermore, their technique is a white-box technique, which extracts features for prediction by analyzing the source code of the program under test, whereas our technique is a black-box technique. That is, with the source code of the program under test, our technique may be further improved to produce better MRs.

8.2 Program-Invariant Inference

As program invariants are important for fault detection and program repairing, researchers proposed to infer program invariants through analysis, especially dynamic analysis. For example, Ernst et al. [21] developed a tool named Daikon to discover program invariants for supporting program evolution. Jiang et al. [32] proposed a novel technique to automatically model and search relationships between the flow intensities that can be regarded as invariants. Csallner et al. [16] proposed to infer invariants using dynamic symbolic execution. Llano et al. [39] proposed to use theory formation to discover invariants. Recently, Nguyen et al. [50] inferred disjunctive invariants with a hybrid approach. Furthermore, as specifications tell the usage of API and may be used to detect faults, many researchers focus

on inferring specifications [81, 82, 73], which can be viewed as program invariants as well. Besides work on invariant generation [79, 37, 6, 49, 42], some researchers focused on using invariants to facilitate software testing (e.g., test-case generation [75] and test-suite reduction [51]), software verification [52], model inference and transformation [35, 9], specifications mining [40], and so on.

Our work is related to program-invariant inference because MRs can also be regarded as program invariants, both of them may be applied to reveal faults in software testing [20]. However, traditional program invariants are supposed to hold during each single execution, whereas MRs are supposed to hold across multiple executions. Similar to dynamic invariant inference, our approach is also based on the analysis of program executions. However, our approach is to search for the values of the parameters in an MR, whereas dynamic variant inference is to discover invariants that satisfying executions.

8.3 Particle Swarm Optimization

Particle swarm optimization (PSO) [54, 80] is a swarm intelligence optimization algorithm simulating the birds foraging behavior. Due to the efficiency of PSO in solving optimization problems, PSO has been applied to various areas, including multi-objective optimization, pattern recognition, signal processing, classification and data clustering.

Recently PSO is applied to some specific areas of software engineering, e.g., automated test-case generation [67]. In this paper, we use PSO for MR inference. To our knowledge, it is the first application of PSO in metamorphic testing.

8.4 Search-Based Software Engineering

Harman and Jones [26] coined the term Search Based Software Engineering (SBSE) and argued that software engineering is ideal for the application of metaheuristic search techniques, such as genetic algorithms. Typically, hill climbing, simulated annealing and genetic algorithms are the three main metaheuristic search techniques that have been widely used in software engineering [25].

Search-based optimization techniques have been widely applied to software testing, including test-suite generation [8, 4, 56, 24, 53] and optimization [72, 31, 43, 3, 74]. Besides software testing, search-based optimization techniques have also been applied to fault localization [65], program analysis [76], software refactoring [29, 30, 55], cost estimation [19], project scheduling [1, 18], decisions design optimization [10], automated negotiation [17], source code parallelization [57], requirement engineering [27, 64], variability management [41], and so on.

Although search-based software engineering is important and promising, very little research in search-based software engineering has used PSO as a metaheuristic search technique. Our work is the first application of search-based software engineering for program-invariant inference.

9. CONCLUSION AND FUTURE WORK

In this paper, we propose a novel approach to automatically inferring polynomial metamorphic relations by analyzing multiple executions of the same program under test. To our knowledge, this is the first automatic approach to MR inference. In particular, we view the problem of MR inference as a searching problem and thus use a typical optimization algorithm PSO to solve the problem. Then we

conducted three empirical studies and got the finding that our approach is able to infer many MRs with high quality in acceptable time, which are effective in detecting faults with no false detection.

In our future, we plan to investigate the following issuers.

First, we will extend types of MRs in future work. Besides polynomial equations studied in this paper, some MRs may be represented by polynomial inequalities. For the scientific function $\log_{10}(x)$, if x_1 is larger than x_2 , $\log_{10}(x_1)$ is larger than $\log_{10}(x_2)$. Relations between programs (e.g., “ $\sin^2(x) + \cos^2(x) = 1$ ”) may also help detect faults in the functions \sin and \cos . In future work, we will extend the definition of MRs and investigate how to infer these MRs.

Second, we will improve our approach by investigating other PSO algorithms or optimization algorithms. Besides the PSO algorithm used in this paper, there exist many other optimization algorithms [1] like hill climbing, which have been used to solve similar search problems (e.g., test-suite reduction) in software testing. In future work, we will investigate some other PSO algorithms for the MR inference problem or optimization algorithms in MR inference.

10. ACKNOWLEDGMENTS

This work is supported by the National Basic Research Program of China under Grant No.2014CB347701, the High-Tech Research and Development Program of China under Grant No.2013AA01A605, and the National Natural Science Foundation of China under Grant Nos.61121063, 61332010, 61272157, 61228203, 61225007.

11. REFERENCES

- [1] E. Alba and F. Chicano. Management of software projects with gas. In *Proc. MIC*, pages 13–18, 2005.
- [2] J. H. Andrews, L. C. Briand, and Y. Labiche. Is mutation an appropriate tool for testing experiments? In *Proc. ICSE*, pages 402–411, 2005.
- [3] J. H. Andrews, T. Menzies, and F. C. Li. Genetic algorithms for randomized unit testing. *IEEE Transactions on Software Engineering*, 37(1):80–94, 2011.
- [4] A. Baars, M. Harman, Y. Hassoun, K. Lakhotia, P. McMinn, P. Tonella, and T. Vos. Symbolic search-based testing. In *Proc. ASE*, pages 53–62, 2011.
- [5] E. T. Barr, T. Vo, V. Le, and Z. Su. Automatic detection of floating-point exceptions. In *Proc. POPL*, pages 549–560, 2013.
- [6] S. Bensalem, M. Bozga, B. Boyer, and A. Legay. Incremental generation of linear invariants for component-based systems. In *Proc. ACSO*, pages 80–89, 2013.
- [7] F. Benz, A. Hildebrandt, and S. Hack. A dynamic program analysis to find floating-point accuracy problems. In *Proc. PLDI*, pages 453–462, 2012.
- [8] A. Bertolino. Software testing research: Achievements, challenges, dreams. In *Proc. FOSE*, pages 85–103, 2007.
- [9] J. Cabot, R. Clarisó, E. Guerra, and J. De Lara. An invariant-based method for the analysis of declarative model-to-model transformations. In *Proc. MODELS*, pages 37–52, 2008.
- [10] G. Canfora and M. Di Penta. New frontiers of reverse engineering. In *Proc. FOSE*, pages 326–341, 2007.

- [11] T. Y. Chen, S. C. Cheung, and S. M. Yiu. Metamorphic testing: A new approach for generating next test cases. Technical Report HKUST-CS98-01, Hong Kong University of Science and Technology, 1998.
- [12] T. Y. Chen, J. Feng, and T. H. Tse. Metamorphic testing of programs on partial differential equations: A case study. In *Proc. COMPSAC*, pages 327–333, 2002.
- [13] T. Y. Chen, D. H. Huang, T. H. Tse, and Z. Q. Zhou. Case studies on the selection of useful relations in metamorphic testing. In *Proc. JIISIC*, pages 569–583, 2004.
- [14] T. Y. Chen, F.-C. Kuo, T. H. Tse, and Z. Q. Zhou. Metamorphic testing and beyond. In *Proc. STEP*, pages 94–100, 2003.
- [15] T. Y. Chen, T. H. Tse, and Z. Q. Zhou. Fault-based testing without the need of oracles. *Information and Software Technology*, 45(1):1–9, 2003.
- [16] C. Csallner, N. Tillmann, and Y. Smaragdakis. Dysy: Dynamic symbolic execution for invariant inference. In *Proc. ICSE*, pages 281–290, 2008.
- [17] E. Di Nitto, M. Di Penta, A. Gambi, G. Ripa, and M. L. Villani. Negotiation of service level agreements: An architecture and a search-based approach. In *Proc. ICSOC*, pages 295–306, 2007.
- [18] M. Di Penta, M. Harman, and G. Antoniol. The use of search-based optimization techniques to schedule and staff software projects: An approach and an empirical study. *Software: Practice and Experience*, 41(5):495–519, 2011.
- [19] J. J. Dolado. A validation of the component-based method for software size estimation. *IEEE Transactions on Software Engineering*, 26(10):1006–1021, 2000.
- [20] M. D. Ernst. *Dynamically discovering likely program invariants*. PhD thesis, University of Washington, 2000.
- [21] M. D. Ernst, J. H. Perkins, P. J. Guo, S. McCamant, C. Pacheco, M. S. Tschantz, and C. Xiao. The Daikon system for dynamic detection of likely invariants. *Science of Computer Programming*, 69(1):35–45, 2007.
- [22] I. M. Gelfand and M. Saul. *Trigonometry*. Springer, 2001.
- [23] P. Godefroid and J. Kinder. Proving memory safety of floating-point computations by combining static and dynamic program analysis. In *Proc. ISSSTA*, pages 1–12, 2010.
- [24] F. Gross, G. Fraser, and A. Zeller. Search-based system testing: high coverage, no false alarms. In *Proc. ISSSTA*, pages 67–77, 2012.
- [25] M. Harman. The current state and future of search based software engineering. In *Proc. FOSE*, pages 342–357, 2007.
- [26] M. Harman and B. F. Jones. Search-based software engineering. *Information and Software Technology*, 43(14):833–839, 2001.
- [27] M. Harman, J. Krinke, J. Ren, and S. Yoo. Search based data sensitivity analysis applied to requirement engineering. In *Proc. GECCO*, pages 1681–1688, 2009.
- [28] M. Harman, P. McMinn, M. Shahbaz, and S. Yoo. A comprehensive survey of trends in oracles for software testing. Technical Report CS-13-01, University of Sheffield, 2013.
- [29] M. Harman and L. Tratt. Pareto optimal search based refactoring at the design level. In *Proc. GECCO*, pages 1106–1113, 2007.
- [30] S. Hayashi, Y. Tsuda, and M. Saeki. Search-based refactoring detection from source code revisions. *IEICE Transactions on Information and Systems*, 93(4):754–762, 2010.
- [31] S. Huang, M. B. Cohen, and A. M. Memon. Repairing GUI test suites using a genetic algorithm. In *Proc. ICST*, pages 245–254, 2010.
- [32] G. Jiang, H. Chen, and K. Yoshihira. Discovering likely invariants of distributed transaction systems for autonomic system management. In *Proc. ICAC*, pages 199–208, 2006.
- [33] U. Kanewala and J. M. Bieman. Using machine learning techniques to detect metamorphic relations for program without test oracles. In *Proc. ISSRE*, pages 1–10, 2013.
- [34] J. Kennedy and R. C. Eberhart. Particle swarm optimization. In *Encyclopedia of Machine Learning*, pages 1942–1948, 1995.
- [35] I. Krka, Y. Brun, D. Popescu, J. Garcia, and N. Medvidovic. Using dynamic execution traces and program invariants to enhance behavioral model inference. In *Proc. ICSE*, pages 179–182, 2010.
- [36] J. J. Liang, A. K. Qin, P. N. Suganthan, and S. Baskar. Comprehensive learning particle swarm optimizer for global optimization of multimodal functions. *IEEE Transactions on Evolutionary Computation*, 10(3):281–295, 2006.
- [37] W. Lin, M. Wu, Z. Yang, and Z. Zeng. Exact safety verification of hybrid systems using sums-of-squares representation. *Science China Information Sciences*, 57(5):1–13, 2014.
- [38] H. Liu, X. Liu, and T. Y. Chen. A new method for constructing metamorphic relations. In *Proc. QSIC*, pages 59–68, 2012.
- [39] M. T. Llano, A. Ireland, and A. Pease. Discovery of invariants through automated theory formation. *Formal Aspects of Computing*, 26(2):203–249, 2014.
- [40] D. Lo and S. Maoz. Mining scenario-based specifications with value-based invariants. In *Proc. OOPSLA*, pages 755–756, 2009.
- [41] R. E. Lopez-Herrejon and A. Egyed. Sbse4vm: Search based software engineering for variability management. In *Proc. CSMR*, pages 441–444, 2013.
- [42] M. Z. Malik, A. Pervaiz, and S. Khurshid. Generating representation invariants of structurally complex data. In *Proc. TACAS*, pages 34–49, 2007.
- [43] N. Mansour, R. Bahsoon, and G. Baradhi. Empirical comparison of regression test selection algorithms. *Journal of Systems and Software*, 57(1):79–90, 2001.
- [44] J. Mayer and R. Guderlei. An empirical study on the selection of good metamorphic relations. In *Proc. COMPSAC*, pages 475–484, 2006.
- [45] J. Mayer and R. Guderlei. An empirical study on the selection of good metamorphic relations. In *Proc. COMPSAC*, pages 475–484, 2006.
- [46] H. Mei, D. Hao, L. Zhang, L. Zhang, J. Zhou, and

- G. Rothermel. A static approach to prioritizing junit test cases. *IEEE Transactions on Software Engineering*, 38(6):1258–1275, 2012.
- [47] C. Murphy. *Metamorphic testing techniques to detect defects in applications without test oracles*. PhD thesis, Columbia University, 2010.
- [48] C. Murphy, K. Shen, and G. Kaiser. Automatic system testing of programs without test oracles. In *Proc. ISSTA*, pages 189–200, 2009.
- [49] T. Nguyen, D. Kapur, W. Weimer, and S. Forrest. Using dynamic analysis to discover polynomial and array invariants. In *Proc. ICSE*, pages 683–693, 2012.
- [50] T. Nguyen, D. Kapur, W. Weimer, and S. Forrest. Using dynamic analysis to generate disjunctive invariants. In *Proc. ICSE*, pages 608–619, 2014.
- [51] N. Pan, F. Zeng, and Y.-H. Huang. Test case reduction based on program invariant and genetic algorithm. In *Proc. WiCOM*, pages 1–5, 2010.
- [52] C. S. Păsăreanu and W. Visser. Verification of java programs using symbolic execution and invariant generation. In *Proc. SPIN*, pages 164–181, 2004.
- [53] Y. Pavlov and G. Fraser. Semi-automatic search-based test generation. In *Proc. ICST*, pages 777–784, 2012.
- [54] R. Poli, J. Kennedy, and T. Blackwell. Particle swarm optimization. *Swarm Intelligence*, 1(1):33–57, 2007.
- [55] F. Qayum and R. Heckel. Search-based refactoring using unfolding of graph transformation systems. *Electronic Communications of the EASST*, 38, 2011.
- [56] D. Romano, M. Di Penta, and G. Antoniol. An approach for search based testing of null pointer exceptions. In *Proc. ICST*, pages 160–169, 2011.
- [57] C. Ryan. *Automatic re-engineering of software using genetic programming*, volume 2. Springer, 2000.
- [58] S. Segura, R. M. Hierons, D. Benavides, and A. Ruiz-Cortés. Automated test data generation on the analyses of feature models: a metamorphic testing approach. In *Proc. ICST*, pages 35–44, 2010.
- [59] S. Segura, R. M. Hierons, D. Benavides, and A. Ruiz-Cortés. Automated metamorphic testing on the analyses of feature models. *Information and Software Technology*, 53(3):245–258, 2011.
- [60] Y. Shi and R. C. Eberhart. Empirical study of particle swarm optimization. In *Proc. CEC*, pages 1945–1950, 1999.
- [61] Y. Shi, H. Liu, L. Gao, and G. Zhang. Cellular particle swarm optimization. *Information Sciences*, 181(20):4460–4493, 2011.
- [62] B. H. Smith and L. Williams. On guiding the augmentation of an automated test suite via mutation analysis. *Empirical Software Engineering*, 14(3):341–369, 2009.
- [63] E. Tang, E. Barr, X. Li, and Z. Su. Perturbing numerical calculations for statistical analysis of floating-point program (in)stability. In *Proc. ISSTA*, pages 131–142, 2010.
- [64] P. Tonella, A. Susi, and F. Palma. Interactive requirements prioritization using a genetic algorithm. *Information and Software Technology*, 55(1):173–187, 2013.
- [65] S. Wang, D. Lo, L. Jiang, and H. C. Lau. Search-based fault localization. In *Proc. ASE*, pages 556–559, 2011.
- [66] E. J. Weyuker. On testing non-testable programs. *The Computer Journal*, 25(4):465–470, 1982.
- [67] A. Windisch, S. Wappler, and J. Wegener. Applying particle swarm optimization to software testing. In *Proc. GECCO*, pages 1121–1128, 2007.
- [68] W.K.Chan, S.C.Cheung, and K. R.P.H.Leung. Towards a metamorphic testing methodology for service-oriented software applications. In *Proc. QSIC*, pages 470–476, 2005.
- [69] W. E. Wong, J. R. Horgan, S. London, and H. Agrawal. A study of effective regression testing in practice. In *Proc. ISSRE*, pages 264–274, 1997.
- [70] P. Wu. Iterative metamorphic testing. In *Proc. COMPSAC*, pages 19–24, 2005.
- [71] X. Xie, J. Ho, C. Murphy, G. Kaiser, B. Xu, and T. Y. Chen. Application of metamorphic testing to supervised classifiers. In *Proc. QSIC*, pages 135–144, 2009.
- [72] Z. Xu, M. B. Cohen, and G. Rothermel. Factors affecting the use of genetic algorithms in test suite augmentation. In *Proc. GECCO*, pages 1365–1372, 2010.
- [73] J. Yang, D. Evans, D. Bhardwaj, T. Bhat, and M. Das. Perracotta: mining temporal API rules from imperfect traces. In *Proc. of ICSE*, pages 282–291, 2006.
- [74] S. Yoo, M. Harman, and S. Ur. GPGPU test suite minimisation: search based software engineering performance improvement using graphics cards. *Empirical Software Engineering*, 18(3):550–593, 2013.
- [75] Y. Yuan, Z. Fanping, Z. Guanmiao, D. Chaoqiang, and X. Neng. Test case generation based on program invariant and adaptive random algorithm. In *Proc. CSE*, pages 274–282, 2011.
- [76] A. Zeller. Search-based program analysis. In *Proc. SSBSE*, pages 1–4, 2011.
- [77] L. Zhang, D. Hao, L. Zhang, G. Rothermel, and H. Mei. Bridging the gap between the total and additional test-case prioritization strategies. In *Proc. ICSE*, pages 192–201. IEEE, 2013.
- [78] L. Zhang, D. Marinov, L. Zhang, and S. Khurshid. Regression mutation testing. In *Proc. ISSTA*, pages 331–341, 2012.
- [79] L. Zhang, G. Yang, N. Rungta, S. Person, and S. Khurshid. Feedback-driven dynamic invariant discovery. In *Proc. ISSTA*, pages 362–372, 2014.
- [80] J. Zhao, C. Han, and B. Wei. Binary particle swarm optimization with multiple evolutionary strategies. *Science China Information Sciences*, 55(11):2485–2494, 2012.
- [81] H. Zhong, T. Xie, L. Zhang, J. Pei, and H. Mei. MAPO: Mining and recommending API usage patterns. In *Proc. of ECOOP*, pages 318–343, 2009.
- [82] H. Zhong, L. Zhang, T. Xie, and H. Mei. Inferring resource specifications from natural language API documentation. In *Proc. of ASE*, pages 307–318, 2009.
- [83] Z. Q. Zhou, D. H. Huang, T. H. Tse, Z. Yang, H. Huang, and T. Y. Chen. Metamorphic testing and its applications. In *Proc. ISFST*, pages 346–351, 2004.